

SMDE FIRST ASSIGMENT

Balbina Virgili Rocosa

27th October 2017

FIRST QUESTION: GENERATE A RANDOM SAMPLE

To start working with probability distributions, first of all, 220 observations have been generated using Spreadsheet.

To do it, next steps have been followed:

1-. 220 random number have been generated

`=RAND()`

2-. The normal distribution of each number generated have been calculated, with mean 0 and standard deviation 1.

`=NORM.INV(AX;0;1)`

3-. The resultant values of the last calculation have been saved on a new document (values_generated.csv).

4-. The observations calculated have been loaded to RStudio

```
#Load the values generated on the spreadsheet file
genDatasetImported <- read.table("./generated_values.csv", header=FALSE)
#The name of the column has been changed just to make it equal as the other subset
colnames(genDatasetImported) <- c("x1")
```

Afterwards, a new collection of 220 observations has been created now with R (with same mean and standard distribution as before).

```
#Generation of a normal distribution
genValuesR <- rnorm(220, mean=0, sd=1)
#Work with the data as a dataframe
genDatasetR <- data.frame(x1=genValuesR)
```

Once both samples are ready to be used, the fitting test of the values is performed in order to determine if exists a statistical dependence between the generated distributions or not. The Chi-squared test of independence will be used for it.

To be able to execute the test, the different values need to be classified into defined intervals.

```
#Definition of the intervals, categories to be used.
tableDatasetImported = transform( genDatasetImported, cat = ifelse(x1 < -1,"-1",
                                                                    ifelse(x1 < -0.5,"-0.5",
                                                                    ifelse(x1 < 0,"0",
                                                                    ifelse(x1 < 0.5,"0.5",
                                                                    ifelse(x1 < 1,"1","Inf"))))))

#Definition of the intervals, categories to be used.
tableDatasetR = transform( genDatasetR, cat = ifelse(x1 < -1,"-1",
                                                                    ifelse(x1 < -0.5,"-0.5",
                                                                    ifelse(x1 < 0,"0",
                                                                    ifelse(x1 < 0.5,"0.5",
                                                                    ifelse(x1 < 1,"1","Inf"))))))
```

The results obtained from the previous classification will be checked in order to see the amount of elements on each frequency.

```
#Counting the amount of elements in each category "table" function.
freqTableImp = as.data.frame(with(tableDatasetImported, table(cat)))
```

```
##      cat Freq
## 1 -0.5   46
## 2  -1   32
## 3   0   42
## 4  0.5   46
## 5   1   33
## 6  Inf   21
```

```
freqTableR = as.data.frame(with(tableDatasetR, table(cat)))
```

```
##      cat Freq
## 1 -0.5   37
## 2  -1   30
## 3   0   45
## 4  0.5   44
## 5   1   29
## 6  Inf   35
```

Both generated frequencies are joined in a new table where the data is classified appropriately to perform the Chi-squared test.

```
#Join both calculated frequencies tables into a final one
freqTableFinal = data.frame(x1 = freqTableImp[2], x2 = freqTableR[2])
```

```
##      Freq Freq.1
## 1    46    37
## 2    32    30
## 3    42    45
## 4    46    44
## 5    33    29
## 6    21    35
```

Finally, the Chi-square test has been performed.

```
#Chi-sqaure test
chiSquareTest = chisq.test(freqTableFinal, correct=FALSE)
```

```
##
## Pearson's Chi-squared test
##
## data:  freqTableFinal
## X-squared = 4.9464, df = 5, p-value = 0.4225
```

With the results obtained, we can appreciate that the p-value calculated is greater than 0.05. Then, the null hypothesis can be accepted and it can be concluded that there is relationship between the two distributions. So, they are not independent and have a similarity. Unfortunately, the level of similarity cannot be calculated with this test.

Now, two other distributions will be analyzed. In this case, two distributions with different mean will be generated and analyzed.

```
#Generate other distributions
genValuesR2 <- rnorm(220, mean=3, sd=3)
genValuesR3 <- rnorm(220, mean=10, sd=4)
```

```

genDatasetR2 <- data.frame(x1=genValuesR2)
genDatasetR3 <- data.frame(x1=genValuesR3)

tableDatasetR2 = transform( genDatasetR2, cat = ifelse(x1 < -1,"-1",
                                                    ifelse(x1 < -0.5,"-0.5",
                                                    ifelse(x1 < 0,"0",
                                                    ifelse(x1 < 0.5,"0.5",
                                                    ifelse(x1 < 1,"1","Inf"))))))))
tableDatasetR3 = transform( genDatasetR3, cat = ifelse(x1 < -1,"-1",
                                                    ifelse(x1 < -0.5,"-0.5",
                                                    ifelse(x1 < 0,"0",
                                                    ifelse(x1 < 0.5,"0.5",
                                                    ifelse(x1 < 1,"1","Inf"))))))))

freqTableR2 = as.data.frame(with(tableDatasetR2, table(cat)))
freqTableR3 = as.data.frame(with(tableDatasetR3, table(cat)))

freqTableFinal2 = data.frame(x1 = freqTableR2[2], x2 = freqTableR3[2])

chiSquareTest2 = chisq.test(freqTableFinal2)

```

```

##
##  Pearson's Chi-squared test
##
## data:  freqTableFinal2
## X-squared = 203.18, df = 5, p-value < 2.2e-16

```

With the results obtained, we can appreciate that the p-value calculated is lower than 0.05. Then, the null hypothesis cannot be accepted and it can be concluded that there is no relationship between the two distributions. So, they are independent and have a no similarity.

SECOND QUESTION: ANOVA

In this exercise, the ANOVA test will be used to test whether three populations have similarity or not.

To start with it, three different distributions have been generated. All of them have different mean values.

#Generation of Three different populations with different mean values

```

v1 = rnorm(200, mean=2, sd=1)
v2 = rnorm(200, mean=5, sd=1)
v3 = rnorm(200, mean=7, sd=1)

```

#Work with the data as a dataframe

```

population1 = data.frame(x1=v1, x2="v1")
population2 = data.frame(x1=v2, x2="v2")
population3 = data.frame(x1=v3, x2="v3")

```

ANOVA test is useful to compare the amount of variation among more than two groups in order to determine if exists a statistical dependence between all of them or not.

The three generated frequencies are joined in a new table where the data is classified appropriately to perform the ANOVA test. In the generated table the first column is representing the calculated value and the second one the category.

```
#RcmdrMisc needs to be loaded to execute the following steps
library("RcmdrMisc")
```

```
## Loading required package: car
```

```
## Loading required package: sandwich
```

```
#Merge all generated distributions into a common table
```

```
data = mergeRows(population1, population2, common.only=FALSE)
```

```
data = mergeRows(as.data.frame(data), population3, common.only=FALSE)
```

Finally, the ANOVA test has been performed.

```
#ANOVA test
```

```
anovaModel <- aov(x1 ~ x2, data=data)
```


```
##              Df Sum Sq Mean Sq F value Pr(>F)
## x2              2  2544.5    1272    1307 <2e-16 ***
## Residuals      597    580.9         1
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

With the results obtained, we can appreciate that the p-value calculated is lower than 0.05. As the p-value is less than the significance level 0.05, we can conclude that there are significant differences between the three distributed models generated.

Unfortunately, ANOVA test does not allow to check if just some of the distributions are similar to others or not. To see it, different pair test would be performed separately.

Below this lines, a visual plot is showed to easily appreciate the results just calculated. With the graphic visualization can be proved that all three distributions are not similar to each other because their mean values are much different.

```
Boxplot(x1~x2, data=data, id.method="y")
```



```
## [1] "300" "398" "430" "453" "551"
```

To prove the test validity, ANOVA test assumptions must be tested.

1-. Normality test

It is necessary to test that the response is normally distributed with Shapiro-Wilk normality test.

```
#Shapiro-Wilk normality test  
shapiro.test(residuals(anovaModel))
```

```
##  
## Shapiro-Wilk normality test  
##  
## data: residuals(anovaModel)  
## W = 0.99458, p-value = 0.03191
```

The calculated p-value is higher than the significance level 0.05, so there is no indication that normality distribution is violated.

2-. Homogeneity of variance

It is necessary to test that the variance is similar within different groups with Breusch Pagan test.

```
#Breusch Pagan test  
lmtest::bptest(anovaModel)
```

```
##  
## studentized Breusch-Pagan test  
##  
## data: anovaModel  
## BP = 0.43991, df = 2, p-value = 0.8026
```

The calculated p-value is higher than the significance level 0.05, so there is no indication of heteroscedasticity on the calculated model. So we can conclude that there is homogeneity of variance. Otherwise, some transformations to eliminate heteroscedasticity would have been performed.

3-. Independence of variables

It is necessary to test that the data values are independent with Durbin Watson test.

```
#Durbin Watson test  
lmtest::dwtest(anovaModel)
```

```
##  
## Durbin-Watson test  
##  
## data: anovaModel  
## DW = 2.0387, p-value = 0.6529  
## alternative hypothesis: true autocorrelation is greater than 0
```

The calculated p-value is higher than the significance level 0, so the alternative hypothesis can be accepted. It means that the variables are not independent and it exists a correlation between them.

Now, a new ANOVA test with a different data will be performed. The data that will be used is geomorphology from FactoMineR package. In fact, the main objective of this part of the assignment is to analyze if drift is a factor that defines different block sizes or wind effects.

First, the data must be loaded.

```
#FactoMineR needs to be loaded to be able to load the desired data  
library("FactoMineR")
```

```
#Load geomorphology dataset
data(geomorphology)
```

Once the data is loaded, the ANOVA test can be executed. To determine if drift is a factor that defines block sizes or wind effects two different ANOVA tests are performed to compare each desired factor with drift one.

```
#ANOVA test to compare block sizes to drifts
anovaModelBlockDrift <- aov(Block.size.median ~ Drift, data=geomorphology)
summary(anovaModelBlockDrift)
```

```
##              Df Sum Sq Mean Sq F value Pr(>F)
## Drift         5    644   128.9    0.726  0.606
## Residuals    69  12246   177.5
```

With the results obtained, we can appreciate that the p-value calculated is higher than 0.05. We can conclude that there are not significant differences between both distributions so we can confirm that drift is a factor that defines different block of sizes.

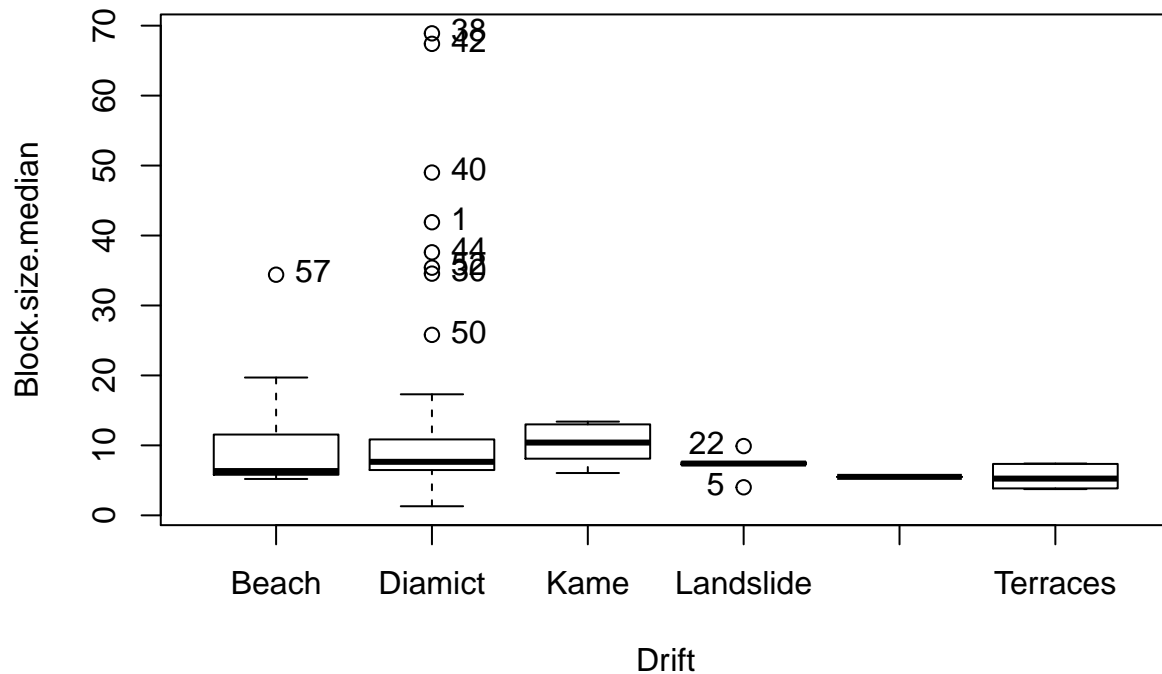
```
#ANOVA test to compare wind effects to drifts
anovaModelWindDrift <- aov(Wind.effect ~ Drift, data=geomorphology)
summary(anovaModelWindDrift)
```

```
##              Df Sum Sq Mean Sq F value    Pr(>F)
## Drift         5  0.2395  0.04790    10.07 2.86e-07 ***
## Residuals    69  0.3280  0.00475
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

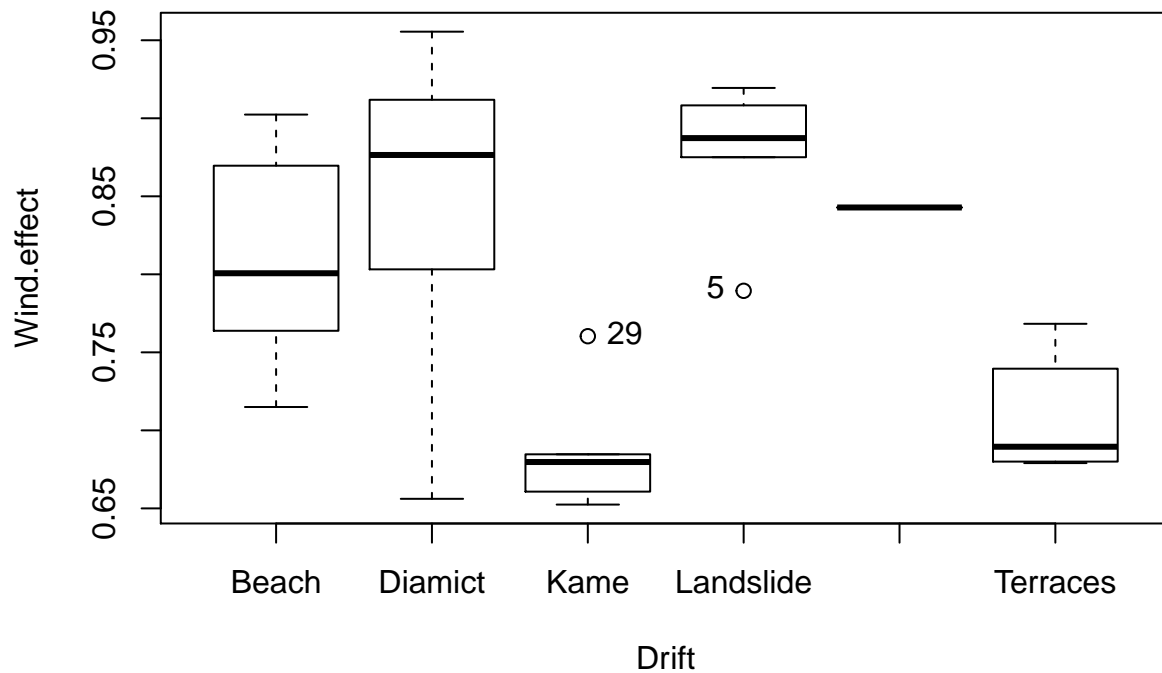
With the results obtained, we can appreciate that the p-value calculated is lower than 0.05. We can conclude that there are significant differences between both distributions so we can confirm that drift is not a factor that defines wind effects.

Finally, the plot of each ANOVA test is shown to check that our conclusions are the correct ones from a visual source.

```
#Boxplot that compare block sizes with drifts
Boxplot(Block.size.median ~ Drift, data=geomorphology, id.method="y")
```



```
## [1] "57" "1" "30" "38" "40" "42" "44" "50" "52" "5" "22"
#Boxplot that compare wind effects with drifts
Boxplot(Wind.effect ~ Drift, data=geomorphology, id.method="y")
```



```
## [1] "29" "5"
```

The first boxplot shows that all different categories have almost the same mean value with a small value of standard variance. While the other boxplot shows the opposite than before, different categories have different means with significant standard variance. So, we may also confirm this way the similarity between the first two factors. Otherwise, the calculation of the ANOVA test is more precise to determine our assumptions.

THIRD QUESTION: DEFINE A LINEAR MODEL FOR AN ATHLETE IN THE 1500 M

In this exercise, the linear expression that better predicts the behavior of an athlete in the 1500m must be determined. The data that will be used is decathlon from FactoMinerR package.

First of all, the data must be loaded.

```
#Load geomorphology dataset  
data(decathlon)
```

Once the data has been loaded, it is time to start playing with it. For it, first test is performed with the all categories that loaded data contains.

```
#Generate the linear model with all data loaded  
lRegression <- lm(`1500m` ~ ., data=decathlon)  
  
##  
## Call:  
## lm(formula = `1500m` ~ ., data = decathlon)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -1.11082 -0.40320 -0.02978  0.37054  0.99005   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)    1.411e+03  2.281e+01  61.851  <2e-16 ***  
## `100m`         -3.711e+01  6.090e-01 -60.937  <2e-16 ***  
## Long.jump      3.947e+01  6.635e-01  59.489  <2e-16 ***  
## Shot.put       9.984e+00  2.172e-01  45.972  <2e-16 ***  
## High.jump     1.470e+02  2.538e+00  57.931  <2e-16 ***  
## `400m`        -7.680e+00  2.351e-01 -32.670  <2e-16 ***  
## `110m.hurdle` -1.951e+01  3.730e-01 -52.320  <2e-16 ***  
## Discus         3.461e+00  5.077e-02  68.184  <2e-16 ***  
## Pole.vault     4.855e+01  6.462e-01  75.127  <2e-16 ***  
## Javeline       2.442e+00  4.648e-02  52.533  <2e-16 ***  
## Rank           3.099e-04  4.278e-02   0.007    0.994  
## Points        -1.632e-01  2.566e-03 -63.589  <2e-16 ***  
## CompetitionOlympicG 1.387e-01  5.532e-01   0.251    0.804  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 0.5796 on 28 degrees of freedom  
## Multiple R-squared:  0.9983, Adjusted R-squared:  0.9975  
## F-statistic: 1350 on 12 and 28 DF, p-value: < 2.2e-16
```

With the results retrieved, we can see that almost all values, except Rank and Competition, have a p-value lower than 0.05 and the t-value much higher or much lower than 0. So we can determine that all categories, except Rank and Competition, have a relation with the 1500m distribution.

Otherwise, the results show that Rank and Competition are not much related with 1500m distribution so both of them can be discarded from the linear model.

Now, a new linear regression model is calculated discarding Rank and Competition distributions.

```
#Generate the linear model without discarded distributions  
lRegression2 <- update(lRegression, . ~ . - Rank - Competition)
```



```
##
## Call:
## lm(formula = `1500m` ~ `100m` + Long.jump + Shot.put + High.jump +
##     `400m` + `110m.hurdle` + Discus + Pole.vault + Javeline +
##     Points, data = decathlon)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.04641 -0.38236 -0.07186  0.39878  0.94814
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.409e+03  1.788e+01   78.81  <2e-16 ***
## `100m`        -3.721e+01  5.541e-01  -67.15  <2e-16 ***
## Long.jump      3.934e+01  5.699e-01   69.03  <2e-16 ***
## Shot.put       9.983e+00  2.074e-01   48.13  <2e-16 ***
## High.jump      1.464e+02  2.090e+00   70.08  <2e-16 ***
## `400m`       -7.627e+00  1.926e-01  -39.60  <2e-16 ***
## `110m.hurdle` -1.947e+01  3.382e-01  -57.57  <2e-16 ***
## Discus         3.451e+00  4.311e-02   80.05  <2e-16 ***
## Pole.vault     4.839e+01  5.317e-01   91.00  <2e-16 ***
## Javeline       2.434e+00  3.966e-02   61.38  <2e-16 ***
## Points        -1.627e-01  1.789e-03  -90.99  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5623 on 30 degrees of freedom
## Multiple R-squared:  0.9983, Adjusted R-squared:  0.9977
## F-statistic: 1721 on 10 and 30 DF, p-value: < 2.2e-16
```

With the results retrieved, we can see that all values have a p-value lower than 0.05 and the t-value much higher or much lower than 0. So we could determine that all categories have a relation with the 1500m distribution.

In this case, the R-squared value is 0.9975 so the probability to predict the 1500m distribution in this calculated linear model can be considered as very good one.

Unfortunately, analyzing all different values used, we can determine that Points must not be used to create the linear model because it is a value calculated from the undetermined value. So, it does not make sense to use this category to predict the 1500m value.

So, a new linear regression model is calculated discarding Points distribution.

```
#Generate the linear model without discarted distribution
lRegression3 <- update(lRegression2, . ~ . - Points)
```

```
##
## Call:
## lm(formula = `1500m` ~ `100m` + Long.jump + Shot.put + High.jump +
##     `400m` + `110m.hurdle` + Discus + Pole.vault + Javeline,
##     data = decathlon)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -20.015  -4.692  -1.217   4.821  20.968
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5.089196 144.658704 -0.035 0.972161
## `100m`      -14.026804   8.055565 -1.741 0.091556 .
## Long.jump   4.363645   6.888604  0.633 0.531080
## Shot.put    -0.712440   2.797866 -0.255 0.800685
## High.jump   -6.663147  20.281067 -0.329 0.744712
## `400m`      7.056379   1.721604  4.099 0.000277 ***
## `110m.hurdle` 0.002113   4.287265  0.000 0.999610
## Discus      1.328956   0.593554  2.239 0.032474 *
## Pole.vault  12.480774   5.834140  2.139 0.040398 *
## Javeline    -0.656775   0.335114 -1.960 0.059056 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9.206 on 31 degrees of freedom
## Multiple R-squared:  0.518, Adjusted R-squared:  0.378
## F-statistic: 3.701 on 9 and 31 DF,  p-value: 0.003022
```

With the new results retrieved, we can see that dependencies have changed. Now, we can determine that 400m, Discus, Pole.vault, Discus, Javeline and 400m have a relation with 1500m distribution. The one that seems to be more related to the 1500m distribution is 400m.

In this case, the R-squared value is 0.378 so the probability to predict the 1500m distribution in this calculated linear model has decreased significantly from the previous calculated model.

To determine that our assumptions are correct, a new linear regression model is calculated discarding all distributions that have been considered as not relationated.

```
#Generate the linear model without discarded distributions
lRegression4 <- lm(formula = `1500m` ~ `100m` + `400m` + Pole.vault + Discus +
                    Javeline, data = decathlon)
```

```
##
## Call:
## lm(formula = `1500m` ~ `100m` + `400m` + Pole.vault + Discus +
##     Javeline, data = decathlon)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -20.831  -3.978  -1.043   4.421  20.784
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  40.5049    85.1779   0.476  0.63736
## `100m`      -15.1377     6.3870  -2.370  0.02343 *
## `400m`       6.6108     1.4152   4.671 4.31e-05 ***
## Pole.vault   13.2420     5.0826   2.605  0.01339 *
## Discus       1.2150     0.4304   2.823  0.00779 **
## Javeline    -0.6844     0.2947  -2.323  0.02613 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.761 on 35 degrees of freedom
## Multiple R-squared:  0.5071, Adjusted R-squared:  0.4367
## F-statistic: 7.201 on 5 and 35 DF,  p-value: 0.0001001
```

With the new results retrieved, we can see that dependencies are still the same as the ones calculated on the

previous model. But now, the R-squared value is 0.4367. So the probability to predict the 1500m distribution in this calculated linear model has increased from the value calculated during the previous model, where much more information was used.

Finally, just to take into account all useful possibilities, a new linear regression model is calculated using just the 400m distribution, which seems to be the most related one from the results retrieved on the previous linear models.

```
#Generate the linear model without discarded distributions
lRegression5 <- lm(formula = `1500m` ~ `400m`, data = decathlon)
```

```
##
## Call:
## lm(formula = `1500m` ~ `400m`, data = decathlon)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.0877  -6.9098  -0.7062   4.7360  31.5996
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    74.102     73.424   1.009  0.31909
## `400m`         4.130       1.479   2.792  0.00808 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 10.79 on 39 degrees of freedom
## Multiple R-squared:  0.1666, Adjusted R-squared:  0.1452
## F-statistic: 7.793 on 1 and 39 DF,  p-value: 0.008078
```

With the new results retrieved, we can see that dependencies have changed a little bit and the R-squared value is 0.1452. So the probability to predict the 1500m distribution in this calculated linear model has decreased almost a 20% from the value calculated during the previous model, where much more information was used.

To conclude, it can be determined that the best linear model to predict the behavior of an athlete for 1500m is the one that takes into account all the following variables:

- 400m
- Pole.vault
- Javeline
- Discus
- 100m

Finally, Linear model test assumptions must be calculated to prove the test validity.

1-. Normality test

It is necessary to test that the response is normally distributed with Shapiro-Wilk normality test.

```
#Shapiro-Wilk normality test
shapiro.test(residuals(lRegression4))
```

```
##
## Shapiro-Wilk normality test
##
## data: residuals(lRegression4)
## W = 0.98581, p-value = 0.8808
```

The calculated p-value is higher than the significance level 0.05, so there is no indication that normality distribution is violated.

2-. Homogeneity of variance

It is necessary to test that the variance is similar within different groups with Breusch Pagan test.

```
#Breusch Pagan test
lmtest::bptest(lRegression4)

##
## studentized Breusch-Pagan test
##
## data: lRegression4
## BP = 8.8268, df = 5, p-value = 0.1162
```

The calculated p-value is higher than the significance level 0.05, so there is no indication of heteroscedasticity on the calculated model. So we can conclude that there is homogeneity of variance. Otherwise, some transformations to eliminate heteroscedasticity would have been performed.

3-. Independence of variables

It is necessary to test that the data values are independent with Durbin Watson test.

```
#Durbin Watson test
lmtest::dwtest(lRegression4)

##
## Durbin-Watson test
##
## data: lRegression4
## DW = 1.4115, p-value = 0.01718
## alternative hypothesis: true autocorrelation is greater than 0
```

The calculated p-value is higher than the significance level 0, so the alternative hypothesis can be accepted. It means that the variables are not independent and it exists a correlation between them.

FOURTH QUESTION: USE THE MODEL TO PREDICT THE BEHAVIOR OF AN ATHLETE

In this exercise, the behavior of an athlete wants to be predicted. It can be done thanks to the linear model calculated on the previous exercise. To do it, decathlon data will be tested.

Firstly, the decathlon data is separated by competition into different variables. This way, the initial data will be divided into training and test data. In this prediction the data obtained from Decastar competition is going to be used as training data and the one obtained from OlympicG is the one that will be used to test our prediction results.

```
#Two new subsets are created with the decathlon data separated by competition
dataDecastar <- subset(decathlon, Competition=='Decastar')
dataOlympicG <- subset(decathlon, Competition=='OlympicG')
```

So the linear model calculated on the previous exercise is again executed but just using the training data.

```
#Generate the linear model with training data
lRegression6 <- lm(formula = `1500m` ~ `100m` + `400m` + Pole.vault, data = dataDecastar)

##
## Call:
## lm(formula = `1500m` ~ `100m` + `400m` + Pole.vault, data = dataDecastar)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -12.421  -3.943  -1.780   2.795  17.172
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   381.959    154.656   2.470  0.03559 *
## `100m`        -40.186     11.492  -3.497  0.00676 **
## `400m`         5.476      3.138   1.745  0.11495
## Pole.vault    16.067     10.022   1.603  0.14335
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.667 on 9 degrees of freedom
## Multiple R-squared:  0.6239, Adjusted R-squared:  0.4986
## F-statistic: 4.977 on 3 and 9 DF,  p-value: 0.02638
```

With the new results retrieved of this new linear model, we can see that R-squared value is 0.5464 so the probability of prediction is above 50% which is considered the minimum percentage accepted on athleticism.

Two different predictions can be performed.

1-. The first one is to predict the behavior of the athletes during the next race. 2-. The second one is to predict the behavior of the athletes during the following races.

As the dataset we are using is not very complete because it does not have lots of data, just the first prediction test will be performed. So, the prediction interval has been calculated using the data sample collected.

#Predict the behavior of the athletes with prediction interval

```
predictedData <- predict(lRegression4, newdata=dataOlympicG, interval="prediction")
```

```
##              fit      lwr      upr
## Sebrle      273.0930 253.0748 293.1111
## Clay        285.7053 265.5114 305.8992
## Karpov       276.6603 256.4152 296.9055
## Macey        276.3674 257.7128 295.0220
## Warners      276.9651 258.2059 295.7244
## Zsivoczky    276.1599 257.8867 294.4330
## Hernu        274.9476 256.7853 293.1099
## Nool          280.3077 260.9544 299.6611
## Bernard      278.2743 259.3566 297.1921
## Schwarzl     283.7812 265.3860 302.1765
## Pogorelov    294.3210 275.4616 313.1805
## Schoenbeck   286.5156 268.0967 304.9345
## Barras       269.7074 251.0662 288.3486
## Smith        275.0395 255.8271 294.2520
## Averyanov    284.1941 264.4627 303.9255
## Ojaniemi     272.9313 253.9383 291.9243
## Smirnov      272.4774 254.2211 290.7337
## Qi           274.1188 255.8486 292.3889
## Drews        276.3185 257.3087 295.3283
## Parkhomenko  278.7023 259.7072 297.6974
## Terek        293.7927 274.2821 313.3034
## Gomez        260.5929 241.4373 279.7484
## Turi         285.6955 266.5427 304.8483
## Lorenzo      267.1617 248.4769 285.8465
```

```
## Karlivans    279.1269 260.4148 297.8390
## Korkizoglou 296.2165 276.9115 315.5216
## Uldal       278.1043 259.5713 296.6374
## Casarsa     297.4954 276.8127 318.1782
```

With the predict function we obtain the ranges within the data of the following race must be determined according to our designed model.

So, we are going to check if the predict function is accurated or not. To do it, we are going to determine if the data from OlympicsG competition actually lies between the ranges calculated.

```
#Determine if the data of OlympicsG competition lies between the rangs calculated
#TRUE if fits into the rang, FALSE OTHERWISE
```

```
compData <- c()
predictedDataFrame <- data.frame(x1=predictedData)
dataOlympicGFrame <- data.frame(x1=dataOlympicG)
for(x in 1:nrow(predictedDataFrame)) {
  if( dataOlympicGFrame$x1.1500m[x] >= predictedDataFrame$x1.lwr[x] &&
      dataOlympicGFrame$x1.1500m[x] <= predictedDataFrame$x1.upr[x]){
    compData[x] <- 'TRUE'
  }else{
    compData[x] <- 'FALSE'
  }
}

predictedDataFrame["prediction"] <- compData
```

##	x1.fit	x1.lwr	x1.upr	prediction
## Sebrle	273.0930	253.0748	293.1111	TRUE
## Clay	285.7053	265.5114	305.8992	TRUE
## Karpov	276.6603	256.4152	296.9055	TRUE
## Macey	276.3674	257.7128	295.0220	TRUE
## Warners	276.9651	258.2059	295.7244	TRUE
## Zsivoczky	276.1599	257.8867	294.4330	TRUE
## Hernu	274.9476	256.7853	293.1099	TRUE
## Nool	280.3077	260.9544	299.6611	TRUE
## Bernard	278.2743	259.3566	297.1921	TRUE
## Schwarzl	283.7812	265.3860	302.1765	TRUE
## Pogorelov	294.3210	275.4616	313.1805	TRUE
## Schoenbeck	286.5156	268.0967	304.9345	TRUE
## Barras	269.7074	251.0662	288.3486	TRUE
## Smith	275.0395	255.8271	294.2520	TRUE
## Averyanov	284.1941	264.4627	303.9255	TRUE
## Ojaniemi	272.9313	253.9383	291.9243	TRUE
## Smirnov	272.4774	254.2211	290.7337	TRUE
## Qi	274.1188	255.8486	292.3889	TRUE
## Drews	276.3185	257.3087	295.3283	TRUE
## Parkhomenko	278.7023	259.7072	297.6974	TRUE
## Terek	293.7927	274.2821	313.3034	TRUE
## Gomez	260.5929	241.4373	279.7484	TRUE
## Turi	285.6955	266.5427	304.8483	TRUE
## Lorenzo	267.1617	248.4769	285.8465	TRUE
## Karlivans	279.1269	260.4148	297.8390	TRUE
## Korkizoglou	296.2165	276.9115	315.5216	FALSE
## Uldal	278.1043	259.5713	296.6374	TRUE
## Casarsa	297.4954	276.8127	318.1782	TRUE

As it was expected, we have obtained 27 well-predicted values and 1 wrong-predicted one. Then, the prediction interval obtained is 96'4% and can be confirmed that the values obtained lie within the prediction interval of 95% of the samples so it can be determined that the model designed during the previous exercise is an accurate one.

FIVE QUESTION: PCA

In this exercise, the PCA analysis will be used to describe the main variables that exists on the dataset decathlon, which has been already used during the previous exercises.


The PCA analysis is useful if to see if there are linear relationships between variables. So, for the analysis, all variables will be used as active variables except Points and Rank which will be used as supplementary variables.

Looking to the loaded data decathlon, Points and Rank variables can be determined to derived data because it can be calculated from the combination of other variables. So, we will use it as supplementary just to see them on the result of the analysis but they are not influencing the active variables initial relationship.


```
#PCA analysis
```

```
pcaAnalysis = PCA(decathlon[,1:12], scale.unit=TRUE, quanti.sup=c(11: 12))
```

Individuals factor map (PCA)



Variables factor map (PCA)



With the results obtained, it can be determined all the relationships between the different characteristics that an athlete is likely to have.

Analyzing the data obtained, it can be assumed that the most linked analysis to the number of pints are the variables 100m, 110m.hurdle, 400m and long jump. While not for 1500m. In fact, 1500m will depend on other values, which were already calculated on the thirtieth exercise, such as Discuss, 400m, Pole.vault...

Additionally, analyzing the map, it can be easily seen that as much faster as an athlete can run 100m, less it will long jump or, viceversa. So, opposite arrows means opposite result of the behavior.

We can also assume that the first two dimensions resume the 50% of the total variance of the dataset.